

# Anomaly detection with Prometheus



# 127.1

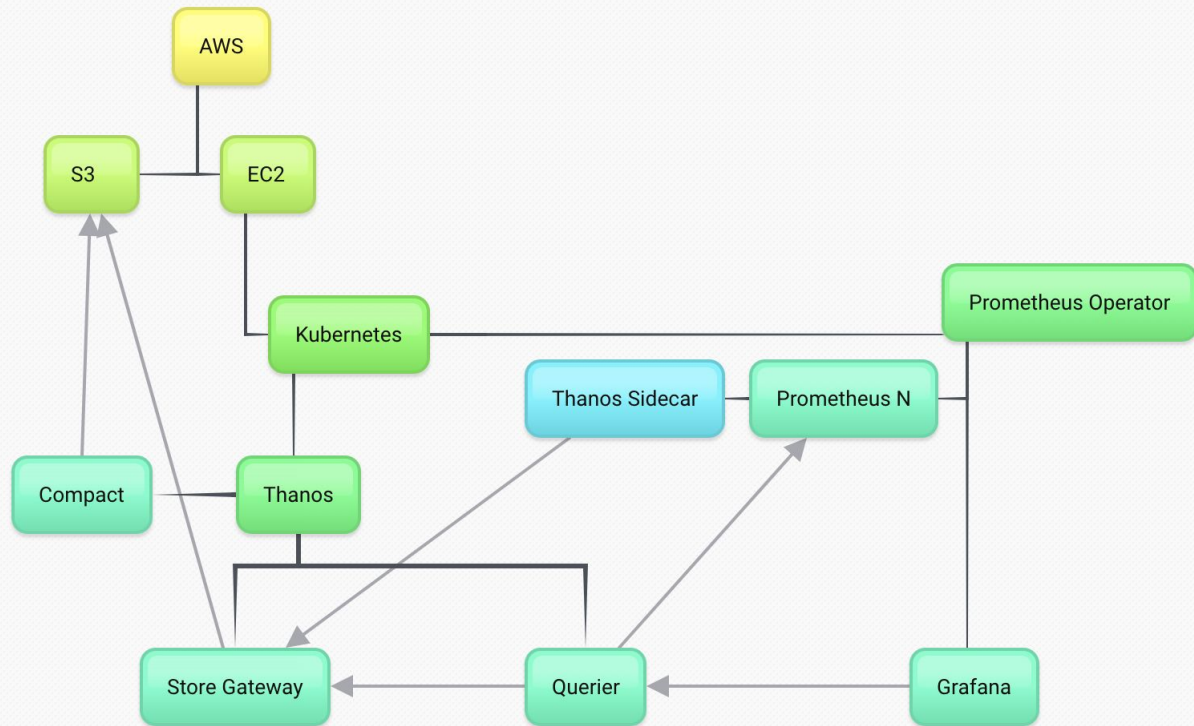
## Jorge Arco

Lead SysAdmin @ [Kodify.io](https://kodify.io)

[github.com/jorge07](https://github.com/jorge07)

# Stack

- Kubernetes in EC2
- Prometheus Operator
- Thanos
- Grafana



# Context

---

Adult industry

High volume

Traffic seasonality

Important number of attacks per week

React faster to possible mistakes



Deeper.COM

**PornTube**<sup>®</sup>

**TUSHY** BLACKED

**VIXEN** BLACKED RAW


# Problem

— — —

“We want to automatically detect  
anomalies in our traffic”

# Understanding the problem

— — —



Those are NOT anomalies

Traffic fluctuates over the year

Marketing campaigns are heavy and can trigger a low traffic alert when they finish

We may be comparing compare against prev months

# Addressing the problem

# Z-Score

— — —

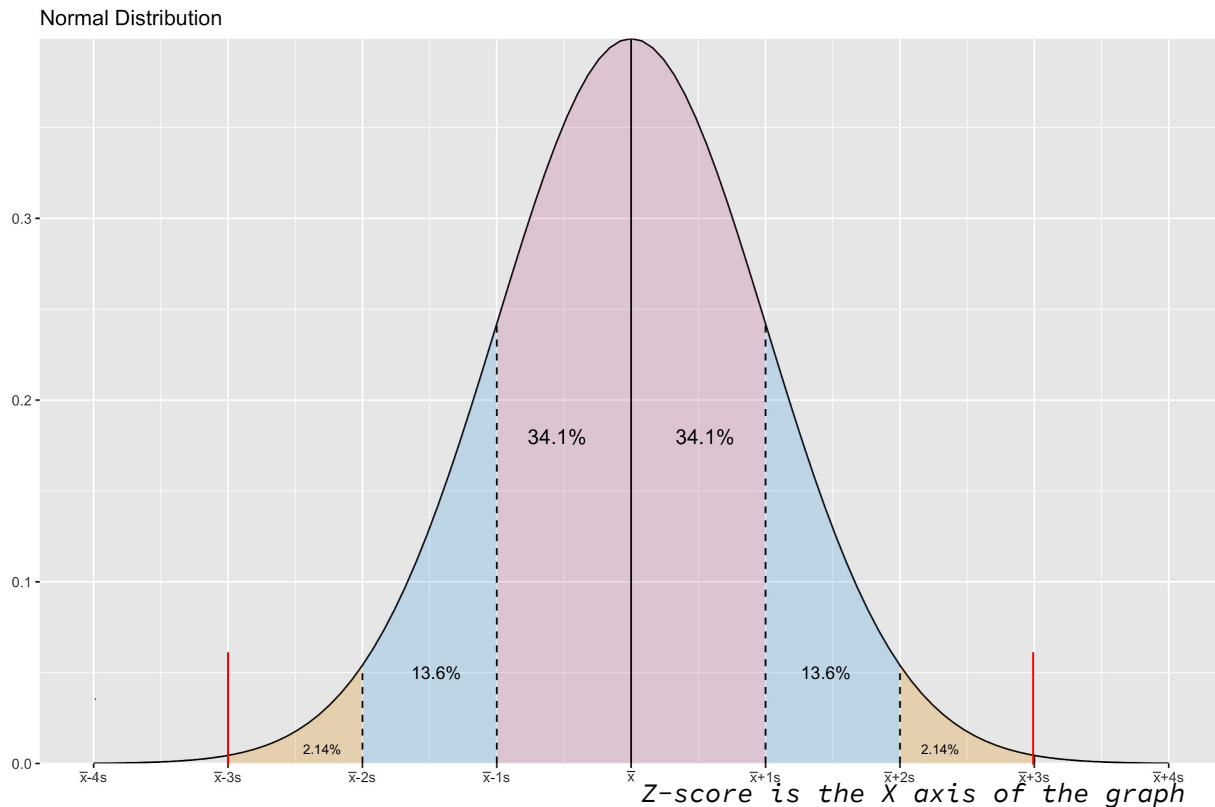
The z-score is measured in the number of standard deviations from the mean

We can assume that any value that falls outside of the range of roughly **+3 to -3** is an **anomaly**

$$z = \frac{x - \mu}{\sigma}$$

$\mu$  = Mean

$\sigma$  = Standard Deviation





# Aggregation

---

- Long term metrics require proper aggregation to avoid expensive queries, use less memory and waste less storage.
- We use **traefik as Ingress controller** so we're going to use this metrics to monitor our traffic

For this example we're going to **exclude all the labels except backend** and aggregate 5 minutes


(backend = www.blacked.com | members.blacked.com)



```
record: traffic:total:sum_rate_5m  
expr: sum by(backend) (rate(traefik_backend_requests_total{job="traefik-frontend-prometheus"}[5m]))
```

# Z-Score calculation in prometheus

— — —



```
- record: traffic:total:sum_rate_5m
  expr: sum by(backend) (rate(traefik_backend_requests_total{job="traefik-frontend-prometheus"}[5m]))

- record: traffic:total:stddev_over_time_1w
  expr: stddev_over_time(traffic:total:sum_rate_5m[1w])

- record: traffic:total:avg_over_time_1w
  expr: avg_over_time(traffic:total:sum_rate_5m[1w])

- record: traffic:total:zscore_1w
  expr: (traffic:total:sum_rate_5m-traffic:total:avg_over_time_1w)/traffic:total:stddev_over_time_1w
```

# Z-score results

---

Sum, week avg and week stddev



## Z-Score



# Seasonality

— — —

## Growth trend per week:

Subtracting the rolling one-week average for last week from the rolling one-week average for now. Take 5 minutes frames from this week and the previous one to generate a prediction. You'll need to play with higher time frames depending of your normal distribution curve inclination (*How variable is you is your traffic timeframe*).

```
- record: traffic:total:sum_rate_5m_prediction_1w
expr: >
  label_replace(
    traffic:total:sum_rate_5m{backend="www.blacked.com/"} offset 1w +
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} -
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} offset 1w
    , "offset", "1w", "", "" )


- record: traffic:total:sum_rate_5m_prediction_2w
expr: >
  label_replace(
    traffic:total:sum_rate_5m{backend="www.blacked.com/"} offset 2w +
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} -
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} offset 2w
    , "offset", "2w", "", "" )

- record: traffic:total:sum_rate_5m_prediction_3w
expr: >
  label_replace(
    traffic:total:sum_rate_5m{backend="www.blacked.com/"} offset 3w +
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} -
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} offset 3w
    , "offset", "3w", "", "" )

- record: traffic:total:sum_rate_5m_prediction_1h_1w
expr: >
  label_replace(
    avg_over_time(
      traffic:total:sum_rate_5m{backend="www.blacked.com/"}[1h] offset 1w
    ) +
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} -
    traffic:total:avg_over_time_1w{backend="www.blacked.com/"} offset 1w
    , "offset", "1w", "", "" )
```

# Seasonality / Trend

— — —



*1, 2 and 3 week trend*

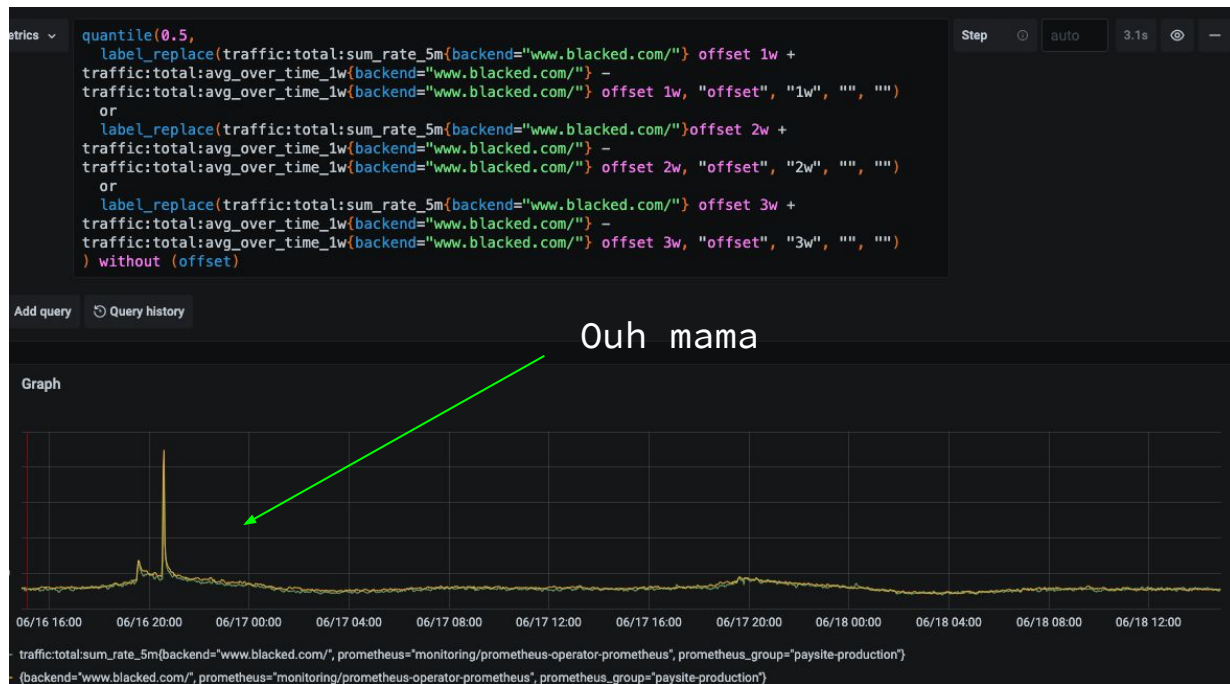
# Seasonality / Trend 1h time span

Let's calculate the median for last 3w and validate results

```
- record: traffic:total:sum_rate_5m_prediction_for_3w
  expr: >
    quantile(0.5,
      traffic:total:sum_rate_5m_prediction_1w
    or
      traffic:total:sum_rate_5m_prediction_2w
    or
      traffic:total:sum_rate_5m_prediction_3w
    ) without (offset)
```



# Seasonality / Trend 5m time span



With the algo validated let's generate the alerts and present this information properly

# Alerting

```
- record: traffic:total:zscore_for_3w_prediction
  expr: >
    (traffic:total:sum_rate_5m - traffic:total:sum_rate_5m_prediction_for_3w )
    / traffic:total:stddev_over_time_1w
```

We're going to take  $\pm 3$   
as our z-score limit


*It doesn't need to be symmetric*

```
- alert: AnomalyInRequestRateInIngress
  expr: >
    abs(
      traffic:total:zscore_for_3w_prediction
    ) > 3
  for: 10m
  labels:
    severity: warning
    team: system
  annotations:
    summary: Requests for Ingress Rule {{ $labels.backend }} are outside normal
```



# Presenting the information

— — —



# Caveats

---

- Long term storage retention in Prometheus cloud be very expensive (Thanos Ruler and Compact components helps a lot here)
- 2-3 Days retention in Prometheus is enough if you move the alert to Ruler component, this saves huge amount of memory
- Ruler component read path is distributed, it depends on network availability and Store Gateway to work, may not fit in all cases.

# Resources

---

[Prometheus](#)

[Thanos - Highly available Prometheus setup with long term storage capabilities](#)

[How to use Prometheus for anomaly detection in GitLab](#)

[Moving Z-Score](#)

QA

— — —

T.Hanks

